

# Duplicate Question Detection Using Machine Learning

## Project Description

This project aims to detect duplicate questions using machine learning techniques. The project processes pairs of questions, extracts various text-based features, and then applies machine learning models to predict whether the questions are duplicates.

## Project Structure

The project consists of three main Jupyter notebooks and a script for training and evaluating models:

### 1. **initial\_EDA.ipynb:**

- Performs initial exploratory data analysis (EDA) on the dataset.
- Analyzes the distribution of duplicate and non-duplicate questions.
- Examines the frequency of unique and repeated questions.

### 2. **only-bow.ipynb:**

- Implements a Bag-of-Words (BoW) approach for feature extraction.
- Merges the text data, applies CountVectorizer to extract features, and splits the data into training and testing sets.

### 3. **bow-with-preprocessing-and-advanced-features.ipynb:**

- Combines advanced preprocessing techniques and feature engineering with BoW.
- Extracts additional features such as the length of the questions, the number of words, common words, and word share between pairs of questions.
- Visualizes these features to understand their distributions and impact on duplicates.
- Applies machine learning models such as Random Forest and XGBoost to classify duplicate questions.

## Usage

### 1. Data Preprocessing:

- Load the dataset using Pandas.
- Remove rows with null values in `question1` and `question2` columns.
- Sample the data for analysis and feature extraction.
- Engineer features like question length, word count, common words, and word share.

### 2. Feature Extraction:

- Apply Bag-of-Words using CountVectorizer to convert text data into numerical features.
- Combine these features with engineered features for model training.

### 3. Model Training:

- Split the data into training and testing sets.
- Train models like Random Forest and XGBoost using the training set.
- Evaluate the models on the test set and measure accuracy.

## Results

- The Random Forest and XGBoost models were trained on the extracted features and achieved the following accuracy scores on the test set:
- **Random Forest:** [insert accuracy]
- **XGBoost:** [insert accuracy]

## Future Work

- Experiment with additional text preprocessing techniques such as TF-IDF and word embeddings.
- Explore deep learning models like LSTM and BERT for improved performance.
- Fine-tune hyperparameters for better model accuracy.

